

Dataset used: Breast Cancer Dataset from the Sklearn library

[illegible]

```

0,      1, 0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 1, 1, 0, 0, 1, 0,
0,      0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 1, 1, 0, 1, 0, 0, 0, 0, 1, 1, 0,
0,      1, 1, 1, 0, 1, 1, 1, 1, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 0, 1,
1,      1, 1, 0, 1, 1, 1, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,      0, 0, 1, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 0, 1, 1, 0, 1, 0, 0, 1,
1,      1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 0, 1, 1, 1, 1,
1,      1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 0, 1, 1, 1, 1, 0,
0,      0, 1, 1, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1,
0,      0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 1, 0, 0, 0, 1, 0,
0,      1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1, 0, 0, 1,
1,      1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 0, 1, 1,
0,      1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 0, 0, 1, 0, 1, 1,
1,      1, 0, 1, 1, 0, 1, 0, 1, 1, 0, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 0,
0,      1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1,
1,      1, 1, 1, 0, 1, 0, 1, 1, 0, 1, 1, 1, 1, 1, 0, 0, 1, 0, 1, 0,
1,      1, 1, 1, 0, 1, 1, 0, 1, 0, 1, 0, 0, 1, 1, 1, 0, 1, 1, 1, 1,
1,      1, 1, 1, 1, 1, 0, 1, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1,      1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 1]),
'frame': None, 'target_names': array(['malignant', 'benign'],
dtype='<U9'), 'DESCR': '.. _breast_cancer_dataset:\n\nBreast cancer
wisconsin (diagnostic) dataset\n
n-----\n\n**Data Set
Characteristics:**\n\nNumber of Instances: 569\n\nNumber of
Attributes: 30 numeric, predictive attributes and the class\n\n
n:Attribute Information:\n      - radius (mean of distances from center
to points on the perimeter)\n      - texture (standard deviation of
gray-scale values)\n      - perimeter\n      - area\n      - smoothness
(local variation in radius lengths)\n      - compactness (perimeter^2 /
area - 1.0)\n      - concavity (severity of concave portions of the
contour)\n      - concave points (number of concave portions of the
contour)\n      - symmetry\n      - fractal dimension ("coastline
approximation" - 1)\n\n      The mean, standard error, and "worst" or

```

```

largest (mean of the three\n    worst/largest values) of these
features were computed for each image,\n    resulting in 30 features.
For instance, field 0 is Mean Radius, field\n    10 is Radius SE,
field 20 is Worst Radius.\n\n    - class:\n        - WDBC-
Malignant\n        - WDBC-Benign\n\nSummary Statistics:\n
n===== \n
Min    Max\n===== \n
radius (mean):                6.981  28.11\ntexture (mean):
9.71  39.28\nperimeter (mean):                43.79  188.5\narea
(mean):                143.5  2501.0\nsmoothness (mean):
0.053  0.163\ncompactness (mean):                0.019  0.345\
nconcavity (mean):                0.0    0.427\nconcave points
(mean):                0.0    0.201\nsymmetry (mean):
0.106  0.304\nfractal dimension (mean):                0.05    0.097\
radius (standard error):                0.112  2.873\ntexture (standard
error):                0.36    4.885\nperimeter (standard error):
0.757  21.98\narea (standard error):                6.802  542.2\
smoothness (standard error):                0.002  0.031\ncompactness
(standard error):                0.002  0.135\nconcavity (standard error):
0.0    0.396\nconcave points (standard error):                0.0    0.053\
symmetry (standard error):                0.008  0.079\nfractal dimension
(standard error):    0.001  0.03\nradius (worst):
7.93  36.04\ntexture (worst):                12.02  49.54\
perimeter (worst):                50.41  251.2\narea (worst):
185.2  4254.0\nsmoothness (worst):                0.071  0.223\
compactness (worst):                0.027  1.058\nconcavity
(worst):                0.0    1.252\nconcave points (worst):
0.0    0.291\nsymmetry (worst):                0.156  0.664\
fractal dimension (worst):                0.055  0.208\
n===== \n\nMissing
Attribute Values: None\n\nClass Distribution: 212 - Malignant, 357 -
Benign\n\nCreator: Dr. William H. Wolberg, W. Nick Street, Olvi L.
Mangasarian\n\nDonor: Nick Street\n\nDate: November, 1995\n\nThis is
a copy of UCI ML Breast Cancer Wisconsin (Diagnostic) datasets.\n
nhttps://goo.gl/U2Uwz2\n\nFeatures are computed from a digitized image
of a fine needle\naspirate (FNA) of a breast mass. They describe\
ncharacteristics of the cell nuclei present in the image.\n\n
nSeparating plane described above was obtained using\nMultisurface
Method-Tree (MSM-T) [K. P. Bennett, "Decision Tree\nConstruction Via
Linear Programming." Proceedings of the 4th\nMidwest Artificial
Intelligence and Cognitive Science Society,\npp. 97-101, 1992], a
classification method which uses linear\nprogramming to construct a
decision tree. Relevant features\nwere selected using an exhaustive
search in the space of 1-4\nfeatures and 1-3 separating planes.\n\nThe
actual linear program used to obtain the separating plane\nin the 3-
dimensional space is that described in:\n[K. P. Bennett and O. L.
Mangasarian: "Robust Linear\nProgramming Discrimination of Two
Linearly Inseparable Sets",\nOptimization Methods and Software 1,
1992, 23-34].\n\nThis database is also available through the UW CS ftp

```

```

server:\n\nftp ftp.cs.wisc.edu\ncd math-prog/cpo-dataset/machine-
learn/WDBC/\n\n.. dropdown:: References\n\n - W.N. Street, W.H.
Wolberg and O.L. Mangasarian. Nuclear feature extraction\n    for
breast tumor diagnosis. IS&T/SPIE 1993 International Symposium on\n
Electronic Imaging: Science and Technology, volume 1905, pages 861-
870,\n    San Jose, CA, 1993.\n - O.L. Mangasarian, W.N. Street and
W.H. Wolberg. Breast cancer diagnosis and\n    prognosis via linear
programming. Operations Research, 43(4), pages 570-577,\n    July-
August 1995.\n - W.H. Wolberg, W.N. Street, and O.L. Mangasarian.
Machine learning techniques\n    to diagnose breast cancer from fine-
needle aspirates. Cancer Letters 77 (1994)\n    163-171.\n',
'feature_names': array(['mean radius', 'mean texture', 'mean
perimeter', 'mean area',
                        'mean smoothness', 'mean compactness', 'mean concavity',
                        'mean concave points', 'mean symmetry', 'mean fractal
dimension',
                        'radius error', 'texture error', 'perimeter error', 'area
error',
                        'smoothness error', 'compactness error', 'concavity error',
                        'concave points error', 'symmetry error',
                        'fractal dimension error', 'worst radius', 'worst texture',
                        'worst perimeter', 'worst area', 'worst smoothness',
                        'worst compactness', 'worst concavity', 'worst concave points',
                        'worst symmetry', 'worst fractal dimension'], dtype='<U23'),
'filename': 'breast_cancer.csv', 'data_module':
'sklearn.datasets.data'}

```

Loading the data into a Pandas dataframe

```

df = pd.DataFrame(breast_cancer_dataset.data, columns =
breast_cancer_dataset.feature_names)

# Distribution of the target variable: Diagnosis

#df['diagnosis'].value_counts() # There is an imbalance in the data,
but as much

# Adding the target column
df['label'] = breast_cancer_dataset.target

```

Exploratory Data Analysis

```

# Printing the first 5 rows of the dataset
df.head()

{"type": "dataframe", "variable_name": "df"}

# Printing the last 5 rows of the dataset
df.tail()

```

```

{"type": "dataframe"}

# Removing the last column as it does not contain any data inplace

#df.drop(columns = "Unnamed: 32", axis = 1, inplace = True)
# To drop a column: Axis = 1
# To drop a row: Axis = 0

df.shape # We have 32 features and 569 samples

(569, 31)

# Removing the column 'id' as it is not necessary for the analysis

#df.drop(columns = "id", axis = 1, inplace = True)

# Encoding the target variable
#label_encoder = LabelEncoder()
#labels = label_encoder.fit_transform(df['diagnosis'])

# Creating a new column for label
#df['target'] = labels

# Dropping the diagnosis column
#df.drop(columns = 'diagnosis', axis = 1, inplace = True)

df.info() # The dataset does not contain any null values

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 31 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   mean radius                           569 non-null    float64
1   mean texture                           569 non-null    float64
2   mean perimeter                         569 non-null    float64
3   mean area                             569 non-null    float64
4   mean smoothness                       569 non-null    float64
5   mean compactness                      569 non-null    float64
6   mean concavity                        569 non-null    float64
7   mean concave points                   569 non-null    float64
8   mean symmetry                         569 non-null    float64
9   mean fractal dimension                569 non-null    float64
10  radius error                          569 non-null    float64
11  texture error                         569 non-null    float64
12  perimeter error                      569 non-null    float64
13  area error                           569 non-null    float64
14  smoothness error                     569 non-null    float64
15  compactness error                    569 non-null    float64
16  concavity error                      569 non-null    float64
17  concave points error                 569 non-null    float64

```

```

18  symmetry error          569 non-null    float64
19  fractal dimension error  569 non-null    float64
20  worst radius            569 non-null    float64
21  worst texture           569 non-null    float64
22  worst perimeter         569 non-null    float64
23  worst area              569 non-null    float64
24  worst smoothness        569 non-null    float64
25  worst compactness       569 non-null    float64
26  worst concavity         569 non-null    float64
27  worst concave points    569 non-null    float64
28  worst symmetry          569 non-null    float64
29  worst fractal dimension  569 non-null    float64
30  label                   569 non-null    int64
dtypes: float64(30), int64(1)
memory usage: 137.9 KB

```

The "label" column is CATEGORICAL as it contains objects which are converted to 0s and 1s to be processed while the other features contain continuous values.

```
df.isnull().sum() # To find if there are any null values
```

```

mean radius          0
mean texture         0
mean perimeter       0
mean area            0
mean smoothness      0
mean compactness     0
mean concavity       0
mean concave points  0
mean symmetry        0
mean fractal dimension 0
radius error         0
texture error        0
perimeter error      0
area error           0
smoothness error     0
compactness error    0
concavity error      0
concave points error 0
symmetry error       0
fractal dimension error 0
worst radius         0
worst texture        0
worst perimeter      0
worst area           0
worst smoothness     0
worst compactness    0
worst concavity      0
worst concave points 0

```

```
worst symmetry          0
worst fractal dimension 0
label                  0
dtype: int64

# Descriptive Statistical measures of the dataset
df.describe() # The function is not carried out for the categorical
feature

{"type": "dataframe"}
```

- 25% of the dataset has a values less than 11.7 in the radius_mean.
- 50% is the median value - middle value after sorting the value

In some cases the mean > median, **right-skewed or positively-skewed data**. And we would need to transform the data into a normally distributed data to make sure that the model makes right predictions.

```
df['label'].value_counts()

# One class has more sample than the other.

label
1      357
0      212
Name: count, dtype: int64

# Finding the mean values of each feature for both the targets.
df.groupby('label').mean()

{"type": "dataframe"}
```

We identify that: 0 -> Benign 1 -> Malignant

The mean values of all the features with the target as Malignant are more than that of the mean with the target as Benign.


Data Visualisation

- Finding the distribution of the features
- Checking for outliers in the dataset
- Correlation between the features

```
import seaborn as sns
import matplotlib.pyplot as plt

# Countplot for the target column
sns.countplot(x = 'label', data = df)

<Axes: xlabel='label', ylabel='count'>
```



Understanding the distribution of features in the dataset

```
# Getting all the column in the dataset  
for column in df:  
    print(column)
```

```
mean radius  
mean texture  
mean perimeter  
mean area  
mean smoothness  
mean compactness  
mean concavity  
mean concave points  
mean symmetry  
mean fractal dimension  
radius error  
texture error  
perimeter error  
area error  
smoothness error  
compactness error  
concavity error  
concave points error  
symmetry error
```





```
fractal dimension error
worst radius
worst texture
worst perimeter
worst area
worst smoothness
worst compactness
worst concavity
worst concave points
worst symmetry
worst fractal dimension
label
```


```
# Univariate analysis - Taking one feature and analysing the  
distribution at a time
```


```
for column in df:
    sns.displot(x = column, data = df)
```


```
/usr/local/lib/python3.10/dist-packages/seaborn/axisgrid.py:453:
RuntimeWarning: More than 20 figures have been opened. Figures created
through the pyplot interface (`matplotlib.pyplot.figure`) are retained
until explicitly closed and may consume too much memory. (To control
this warning, see the rcParam `figure.max_open_warning`). Consider
using `matplotlib.pyplot.close()`.
    fig = plt.figure(figsize=figsize)
```

























































































We can understand the **most of the features are positively skewed**. IE, the mean is greater the median, meaning we have more samples for one condition and less data in the other leading to a case where the model finds it difficult to make predictions for the regions where it has less data in.

```
sns.distplot(x = df['mean radius'])
```

<ipython-input-21-df84fe73072b>:1: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(x = df['mean radius'])
```


<Axes: ylabel='Density'>




```
# Finding correlation between the features - Bi-variate analysis
# sns.pairplot(df)

# Generate 30*30 graphs - Time consuming
# Making a pairplot for 2 features

first_feature = df.iloc[:, 0]
second_feature = df.iloc[:, 1]

plt.scatter(x = first_feature, y = second_feature)
<matplotlib.collections.PathCollection at 0x78dac18127d0>
```



Outlier Detection

An observation that lies an abnormal distance from other values in a random sample from a population.


Regression models are sensitive to outliers and are needed to be processed.



```
# Box and Whisker plot
for column in df:
    plt.figure(figsize = (5,5))
    df.boxplot([column])



# The black circle represents the outliers, with the blue horizontal
# lines representing the 25th(Q1) and 75th(Q3) percentile. The green
# line(Q2) represents the median.
# The black horizontal line are the minimum and maximum values
# Q3 - Q1 is the interquartile range
# Values that lie above 1.5 times the inter-quartile range are the
# outliers.
```



<ipython-input-24-be0ffa52dc02>:3: RuntimeWarning: More than 20 figures have been opened. Figures created through the pyplot interface (`matplotlib.pyplot.figure`) are retained until explicitly closed and may consume too much memory. (To control this warning, see the rcParam `figure.max_open_warning`). Consider using



```
`matplotlib.pyplot.close()`.  
plt.figure(figsize = (5,5))
```
























































```
# When using dataset such as this to model a regression model, the  
outliers need to be replaced.  
# We replace them using the median(more)/mean(less) value.  
  
correlation_matrix = df.corr()  
plt.figure(figsize = (20,20))  
sns.heatmap(correlation_matrix, cbar = True, fmt = '.1f', annot =  
True, cmap = 'crest')  
  
# Saving the image  
plt.savefig('correlation_matrix.png')
```



Observation: The correlation matrix shows the relationships between different features (variables) in a dataset. The diagonal values of the matrices are always 1, as each feature is perfectly correlated with itself.

Correlation Coefficients: Values range between -1 and +1:

+1: Perfect positive correlation (as one feature increases, the other increases proportionally). -1: Perfect negative correlation (as one feature increases, the other decreases proportionally). 0: No correlation (features are independent).

Feature Relationships: Identifying the highly correlated features (positive or negative). Features with high correlation (close to +1 or -1) may be redundant. Detect multicollinearity, where independent variables are highly correlated with one or more variables, which can affect the performance of models like regression.

Feature Selection: Remove redundant features with high correlation to simplify models and improve efficiency

```
# Separating the features and target

X = df.drop(columns = 'label', axis = 1)
y = df['label']

# Splitting data into training and testing data
# 80% - Training data
# 20% - Testing data

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size =
0.2, random_state = 42)

print(X.shape, X_train.shape, X_test.shape)

(569, 30) (455, 30) (114, 30)

print(y_train.shape, y_test.shape)

(455,) (114,)
```

Training a Logistic Regression Model

```
model = LogisticRegression()

# Training the model using Training
model.fit(X_train, y_train)

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs failed to converge
(status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as
shown in:
    https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-
regression
    n_iter_i = _check_optimize_result(
LogisticRegression()
```

```
# Model evaluation - Accuracy on training data
```

```
train_pred = model.predict(X_train)
train_accuracy = accuracy_score(y_train, train_pred)
print("The Training accuracy: ", train_accuracy)
```

```
The Training accuracy:  0.9472527472527472
```

```
y_pred = model.predict(X_test)
test_accuracy = accuracy_score(y_test, y_pred)
print("The test accuracy: ", test_accuracy)
```

```
The test accuracy:  0.956140350877193
```

Predictive System

```
input_data = (19.81, 22.15, 130, 1260, 0.09831, 0.1027, 0.1479,
0.09498, 0.1582, 0.05395, 0.7582, 1.017, 5.865, 112.4, 0.006494,
0.01893, 0.03391, 0.01521, 0.01356, 0.001997, 27.32, 30.88, 186.8,
2398, 0.1512, 0.315, 0.5372, 0.2388, 0.2768, 0.07615)
input_array = np.asarray(input_data)
```

```
# Reshaping the data
```

```
input_array_reshaped = input_array.reshape(1, -1)
```

```
prediction = model.predict(input_array_reshaped)
print(prediction)
```

```
if prediction[0] == 0:
    print("The patient is Benign")
else:
    print("The patient is Malignant")
```

```
[0]
The patient is Benign
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/utils/
validation.py:2739: UserWarning: X does not have valid feature names,
but LogisticRegression was fitted with feature names
  warnings.warn(
```

Building a Neural Network

```
# Standarise the data
```

```
from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()
X_train_standard = scaler.fit_transform(X_train)
X_test_standard = scaler.transform(X_test)
```

```

import tensorflow as tf
# Setting seed
tf.random.set_seed(3) # For reproducibility

from tensorflow import keras

X.shape
(569, 30)

# Setting the layers for the neural network
model = keras.Sequential([
    keras.layers.Flatten(input_shape = (30,)),
    # layer 1 - input layer
    keras.layers.Dense(20, activation =
'relu'), # layer 2 - hidden layer with 20 neurons
    keras.layers.Dense(2, activation =
'sigmoid') # layer 3 - output layer
])

# The input is Flattened to a single dimensional array
# Dense means that all the neurons in the particular layer is
connected to all the neurons in the previous layer.
# Number of neurons in the output layer should be equal to the number
of class in the Target.

# Compiling the neural network
model.compile(optimizer = 'adam',
              loss = 'sparse_categorical_crossentropy', # Label
              encoding
              metrics = ['accuracy']
              )

# Training the neural network

history = model.fit(X_train_standard, y_train, validation_split = 0.1,
epochs = 10)

Epoch 1/10
13/13 _____ 0s 13ms/step - accuracy: 0.5753 - loss:
0.7472 - val_accuracy: 0.8696 - val_loss: 0.3590
Epoch 2/10
13/13 _____ 0s 6ms/step - accuracy: 0.8758 - loss:
0.3308 - val_accuracy: 0.9130 - val_loss: 0.2176
Epoch 3/10
13/13 _____ 0s 10ms/step - accuracy: 0.9180 - loss:
0.2126 - val_accuracy: 0.9348 - val_loss: 0.1796
Epoch 4/10
13/13 _____ 0s 4ms/step - accuracy: 0.9493 - loss:

```

```
0.1670 - val_accuracy: 0.9565 - val_loss: 0.1623
Epoch 5/10
13/13 _____ 0s 5ms/step - accuracy: 0.9520 - loss:
0.1437 - val_accuracy: 0.9565 - val_loss: 0.1514
Epoch 6/10
13/13 _____ 0s 5ms/step - accuracy: 0.9580 - loss:
0.1298 - val_accuracy: 0.9565 - val_loss: 0.1427
Epoch 7/10
13/13 _____ 0s 4ms/step - accuracy: 0.9620 - loss:
0.1199 - val_accuracy: 0.9565 - val_loss: 0.1355
Epoch 8/10
13/13 _____ 0s 4ms/step - accuracy: 0.9683 - loss:
0.1120 - val_accuracy: 0.9565 - val_loss: 0.1292
Epoch 9/10
13/13 _____ 0s 4ms/step - accuracy: 0.9683 - loss:
0.1053 - val_accuracy: 0.9783 - val_loss: 0.1238
Epoch 10/10
13/13 _____ 0s 5ms/step - accuracy: 0.9662 - loss:
0.0994 - val_accuracy: 0.9783 - val_loss: 0.1191
```

We can observe that at each epoch, the accuracy score is increasing and the loss is decreasing.


=> Loss function and accuracy are inversely proportional.

Plotting the accuracy and loss

```
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])

plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')

plt.legend(['Train', 'Validation'], loc = 'upper left')
plt.show()
```





```
# Plotting the accuracy and loss
```

```
plt.plot(history.history['loss'])  
plt.plot(history.history['val_loss'])
```

```
plt.title('Model Accuracy')  
plt.ylabel('Accuracy')  
plt.xlabel('Epoch')
```

```
plt.legend(['Train', 'Validation'], loc = 'upper left')  
plt.show()
```



```
# Accuracy of the model with the test data
```

```
# The model uses the test data to find y_pred(Prediction) and compares it with y_test(True Value) and gives us the loss value and accuracy
```

```
loss, accuracy = model.evaluate(X_test_standard, y_test)
```

```
print(accuracy)
```

```
4/4 ————— 0s 5ms/step - accuracy: 0.9491 - loss: 0.1161
```

```
0.9561403393745422
```

```
print(X_test_standard.shape)
```

```
print("\nThe first sample is:")
```

```
print(X_test_standard[0])
```

```
(114, 30)
```

```
The first sample is:
```

```
[-0.46649743 -0.13728933 -0.44421138 -0.48646498  0.28085007
```

```
0.04160589
```

```
-0.11146496 -0.26486866  0.41524141  0.13513744 -0.02091509 -
```

```
0.29323907
```

```
-0.17460869 -0.2072995  -0.01181432 -0.35108921 -0.1810535 -
```

```
0.24238831
```

```

-0.33731758 -0.0842133 -0.2632354 -0.14784208 -0.33154752 -
0.35109337
 0.48001942 -0.09649594 -0.03583041 -0.19435087  0.17275669
0.20372995]

y_pred = model.predict(X_test_standard) # gives the probability of
each class
print(y_pred.shape)
print("The prediction for the first sampleis:", y_pred[0])

4/4 _____ 0s 4ms/step
(114, 2)
The prediction for the first sampleis: [0.23526107 0.6011053 ]

```

The model has 2 columns, the probabilities of the datapoint being in each of the classes, meaning the model is 23.5% sure that the datapoint is Malignant and 60.1% sure that it is Benign.

```

print(y_pred)

[[0.23526107 0.6011053 ]
 [0.967012   0.17916565]
 [0.8809996  0.39203843]
 [0.09163497 0.76184845]
 [0.0518633  0.87765497]
 [0.997543   0.03363311]
 [0.9975007  0.03328505]
 [0.826613   0.42277163]
 [0.6363133  0.65119433]
 [0.04099468 0.7786242 ]
 [0.08233156 0.5207133 ]
 [0.76948696 0.28281498]
 [0.05472853 0.5836183 ]
 [0.4352791  0.5454692 ]
 [0.05069268 0.83709496]
 [0.5420243  0.16779551]
 [0.04411402 0.7464872 ]
 [0.00502491 0.81901234]
 [0.00185251 0.8099445 ]
 [0.9899725  0.08876199]
 [0.39394748 0.56767327]
 [0.09328163 0.71416384]
 [0.99696106 0.08862436]
 [0.00634144 0.6039927 ]
 [0.03394335 0.79270506]
 [0.08002656 0.82699513]
 [0.06072469 0.8766667 ]
 [0.06778082 0.6532684 ]
 [0.05446557 0.62938887]
 [0.9559362  0.27618167]
 [0.05411643 0.754101  ]

```

[0.01126605 0.7540706]
[0.0531918 0.8202688]
[0.05752247 0.5818418]
[0.01597529 0.7029848]
[0.01669249 0.76145005]
[0.45888856 0.57793117]
[0.03493489 0.82175535]
[0.96165407 0.3150286]
[0.12034945 0.55135024]
[0.02004644 0.8215898]
[0.5882047 0.25820234]
[0.06402855 0.7974846]
[0.02627368 0.68792146]
[0.0835228 0.549306]
[0.10372543 0.6661686]
[0.02972844 0.8801018]
[0.00777749 0.76747096]
[0.15875538 0.6356329]
[0.08850199 0.7609377]
[0.9634779 0.34864497]
[0.9723907 0.06767032]
[0.65766454 0.46065432]
[0.33900797 0.58311003]
[0.04361732 0.7901425]
[0.12744749 0.67763054]
[0.04798314 0.82949865]
[0.99901146 0.00516789]
[0.8659149 0.5371398]
[0.01119028 0.7887057]
[0.06462906 0.621182]
[0.9917601 0.09564156]
[0.9896463 0.0128862]
[0.11215623 0.5670966]
[0.01623746 0.6587871]
[0.3202244 0.6376673]
[0.9205148 0.14695016]
[0.9960741 0.02753112]
[0.01808445 0.760675]
[0.19003838 0.711686]
[0.56790584 0.2588996]
[0.9079595 0.43957266]
[0.06634501 0.56892574]
[0.9246049 0.23878777]
[0.00576476 0.7603435]
[0.1334517 0.6001706]
[0.10732 0.42613506]
[0.6488602 0.6484444]
[0.00486984 0.81886077]
[0.15103422 0.7189332]

```

[0.81270826 0.43475103]
[0.00585179 0.7411847 ]
[0.5792724  0.5233446 ]
[0.9951409  0.03793609]
[0.6241894  0.4972551 ]
[0.9292105  0.13641322]
[0.94384825 0.2213732 ]
[0.8429959  0.08229268]
[0.04146893 0.600876 ]
[0.07614221 0.61682224]
[0.09645507 0.791111 ]
[0.4593728  0.60592794]
[0.21548477 0.3536443 ]
[0.00850809 0.5197519 ]
[0.00946542 0.8385546 ]
[0.01160217 0.84138703]
[0.8301103  0.07711208]
[0.9167709  0.14343035]
[0.00574798 0.8308903 ]
[0.8908912  0.19107935]
[0.8219067  0.55518186]
[0.00107736 0.8272977 ]
[0.9718593  0.10672905]
[0.9428982  0.24071299]
[0.0553957  0.69913006]
[0.13274953 0.46750888]
[0.12862155 0.6439526 ]
[0.9964425  0.03770087]
[0.40135416 0.69930834]
[0.16799852 0.67022586]
[0.8898151  0.18716756]
[0.03737961 0.73771644]
[0.3370565  0.33216566]
[0.96457285 0.0465486 ]]

```

If the first value is greater, then the output is 0. If the second value is greater then, the label is 1.

```
# Converting the probabilities into labels
```

```

y_labels = [np.argmax(i) for i in y_pred]
print(y_labels)

```

```

[1, 0, 0, 1, 1, 0, 0, 0, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1, 0,
1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 1, 1, 1,
1, 1, 1, 1, 0, 0, 0, 1, 1, 1, 1, 0, 0, 1, 1, 0, 0, 1, 1, 1, 0, 0, 1,
1, 0, 0, 1, 0, 1, 1, 1, 0, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1,
1, 1, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 1, 0, 1, 1, 0, 0]

```

Building a Predictive System

```
input_data = (19.81, 22.15, 130, 1260, 0.09831, 0.1027, 0.1479,  
0.09498, 0.1582, 0.05395, 0.7582, 1.017, 5.865, 112.4, 0.006494,  
0.01893, 0.03391, 0.01521, 0.01356, 0.001997, 27.32, 30.88, 186.8,  
2398, 0.1512, 0.315, 0.5372, 0.2388, 0.2768, 0.07615)
```

```
# Converting the input data into numpy array
```

```
input_array = np.asarray(input_data)
```

```
# reshaping the input array
```

```
input_array_reshaped = input_array.reshape(1, -1)
```

```
# Standardise the input data
```

```
input_data_standard = scaler.transform(input_array_reshaped)
```

```
# Making the prediction
```

```
prediction = model.predict(input_data_standard)
```

```
print(prediction)
```

```
label = [np.argmax(prediction)]
```

```
print(label)
```

```
if label[0] == 0:
```

```
    print("The patient is Benign")
```

```
else:
```

```
    print("The patient is Malignant")
```

```
1/1 _____ 0s 29ms/step
```

```
[[0.99331516 0.06749561]]
```

```
[0]
```

```
The patient is Benign
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/utils/  
validation.py:2739: UserWarning: X does not have valid feature names,  
but StandardScaler was fitted with feature names  
warnings.warn(
```